

# Today's Agenda

- Previous class review (refresher)
  - Routing
    - Introduction
    - Distance Vector
    - Link State
- Today's Class
  - UDP Protocol
  - TCP Protocol
    - Connection setup
    - State Transitions
    - Connection Teardown

# Due Dates

- Programming Assignment I (due March 29, 11:59PM)
- ~~Term Paper Email (due February 5, 11:59PM)~~
- ~~Term Paper Proposal (due February 19, 11:59PM)~~
- **Midterm Examination (March 12)**
- Full term Paper (due May 6, 11:59PM)

# Previous Class

# Routing

## Forwarding versus Routing

- Forwarding:
  - to select an output port based on destination address and routing table
- Routing:
  - process by which routing table is built

# Routing

- Forwarding table VS Routing table
  - Forwarding table
    - Used when a packet is being forwarded and so must contain enough information to accomplish the forwarding function
    - A row in the forwarding table contains the mapping from a network number to an outgoing interface and some MAC information, such as Ethernet Address of the next hop
  - Routing table
    - Built by the routing algorithm as a precursor to build the forwarding table
    - Generally contains mapping from network numbers to next hops

# Routing

(a)

Prefix/Length	Next Hop
18/8	171.69.245.10

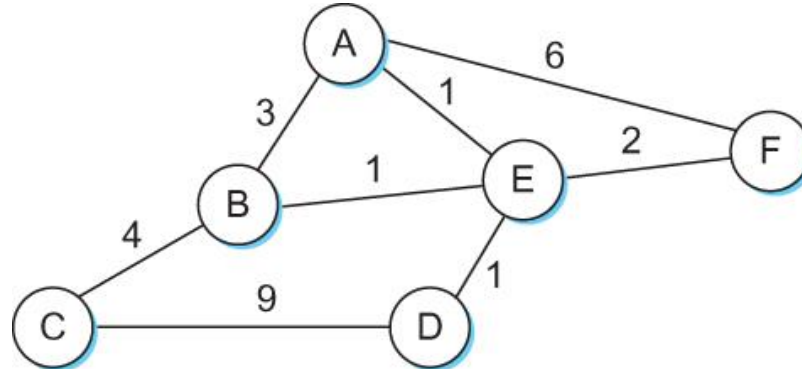
(b)

Prefix/Length	Interface	MAC Address
18/8	if0	8:0:2b:e4:b:1:2

Example rows from (a) routing and (b) forwarding tables

# Routing

- Network as a Graph



- The basic problem of routing is to find the lowest-cost path between any two nodes
  - Where the cost of a path equals the sum of the costs of all the edges that make up the path

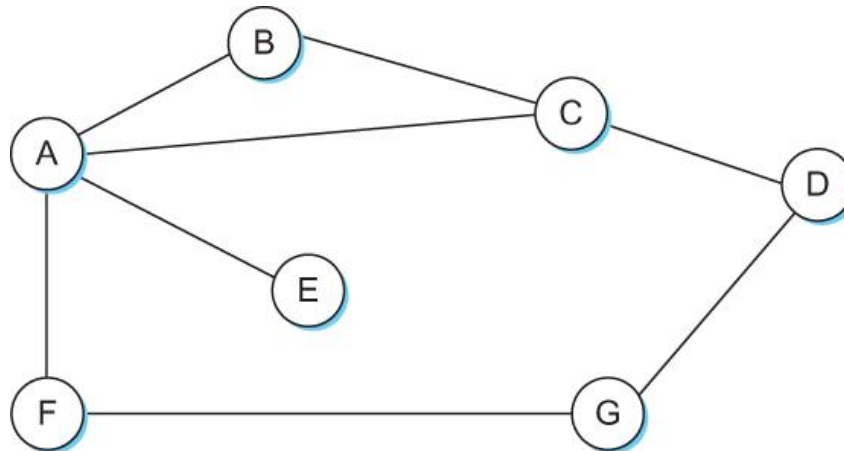
# Routing

- For a simple network, we can calculate all shortest paths and load them into some nonvolatile storage on each node.
- Such a static approach has several shortcomings
  - It does not deal with node or link failures
  - It does not consider the addition of new nodes or links
  - It implies that edge costs cannot change
- What is the solution?
  - Need a distributed and dynamic protocol
  - Two main classes of protocols
    - Distance Vector
    - Link State

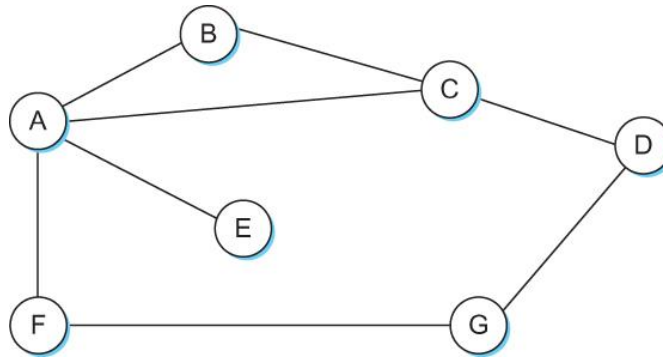


# Distance Vector

- Each node constructs a one dimensional array (a vector) containing the “distances” (costs) to all other nodes and distributes that vector to its immediate neighbors
- Starting assumption is that each node knows the cost of the link to each of its directly connected neighbors



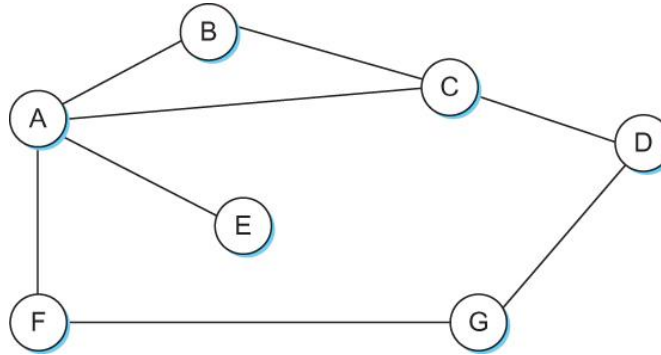
# Distance Vector



Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	$\infty$	1	1	$\infty$
B	1	0	1	$\infty$	$\infty$	$\infty$	$\infty$
C	1	1	0	1	$\infty$	$\infty$	$\infty$
D	$\infty$	$\infty$	1	0	$\infty$	$\infty$	1
E	1	$\infty$	$\infty$	$\infty$	0	$\infty$	$\infty$
F	1	$\infty$	$\infty$	$\infty$	$\infty$	0	1
G	$\infty$	$\infty$	$\infty$	1	$\infty$	1	0

Initial distances stored at each node (global view)

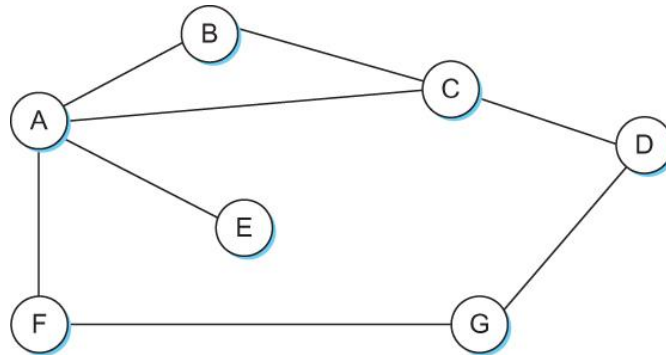
# Distance Vector



Destination	Cost	NextHop
B	1	B
C	1	C
D	$\infty$	—
E	1	E
F	1	F
G	$\infty$	—

Initial routing table at node A

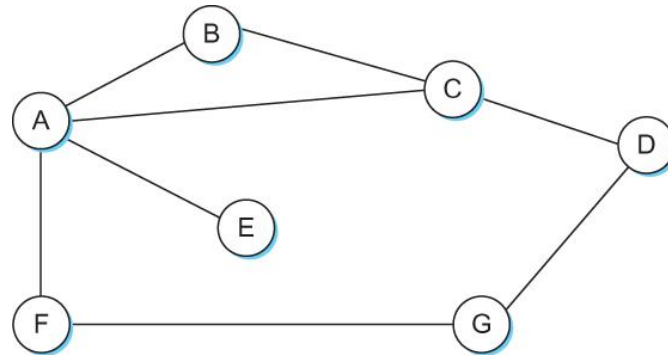
# Distance Vector



Destination	Cost	NextHop
B	1	B
C	1	C
D	2	C
E	1	E
F	1	F
G	2	F

Final routing table at node A

# Distance Vector



Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	3
C	1	1	0	1	2	2	2
D	2	2	1	0	3	2	1
E	1	2	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	3	2	1	3	1	0

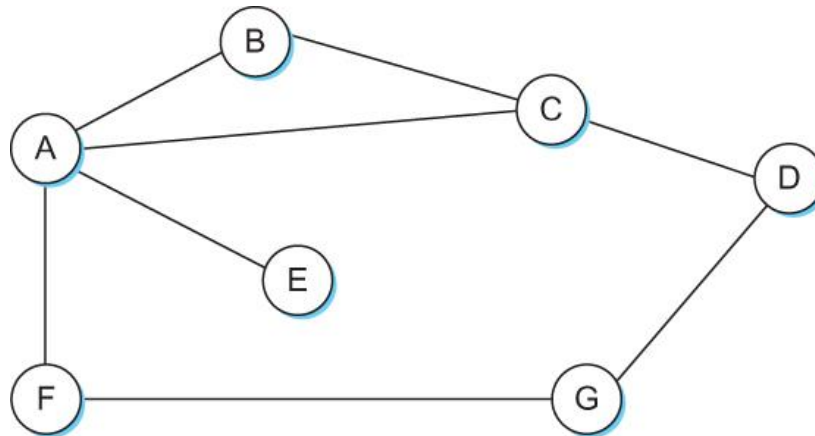
Final distances stored at each node (global view)

# Distance Vector

- The distance vector routing algorithm is sometimes called as Bellman-Ford algorithm
- Every  $T$  seconds each router sends its table to its neighbor each router then updates its table based on the new information
- Problems include fast response to good news and slow response to bad news. Also too many messages to update

# Distance Vector

- When a node detects a link failure
  - F detects that link to G has failed
  - F sets distance to G to infinity and sends update to A
  - A sets distance to G to infinity since it uses F to reach G
  - A receives periodic update from C with 2-hop path to G
  - A sets distance to G to 3 and sends update to F
  - F decides it can reach G in 4 hops via A



# Distance Vector

- Slightly different circumstances can prevent the network from stabilizing
  - Suppose the link from A to E goes down
  - In the next round of updates, A advertises a distance of infinity to E, but B and C advertise a distance of 2 to E
  - Depending on the exact timing of events, the following might happen
    - Node B, upon hearing that E can be reached in 2 hops from C, concludes that it can reach E in 3 hops and advertises this to A
    - Node A concludes that it can reach E in 4 hops and advertises this to C
    - Node C concludes that it can reach E in 5 hops; and so on.
    - This cycle stops only when the distances reach some number that is large enough to be considered infinite
      - **Count-to-infinity problem**



# Count-to-infinity Problem

- Use some relatively small number as an approximation of infinity
- For example, the maximum number of hops to get across a certain network is never going to be more than 16
- One technique to improve the time to stabilize routing is called *split horizon*
  - When a node sends a routing update to its neighbors, it does not send those routes it learned from each neighbor back to that neighbor
  - For example, if B has the route (E, 2, A) in its table, then it knows it must have learned this route from A, and so whenever B sends a routing update to A, it does not include the route (E, 2) in that update

# Count-to-infinity Problem

- In a stronger version of split horizon, called *split horizon with poison reverse*
  - B actually sends that back route to A, but it puts negative information in the route to ensure that A will not eventually use B to get to E
  - For example, B sends the route (E,  $\infty$ ) to A

# **EXERCISE 46, PAGE 294**

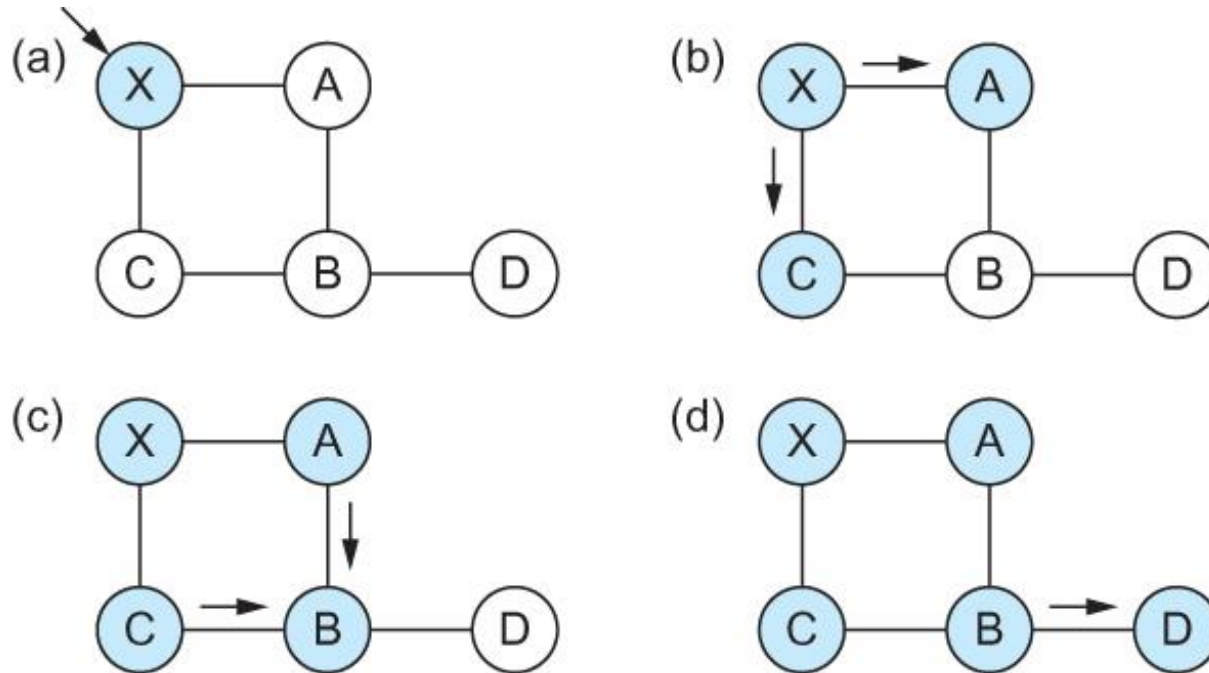
# Link State Routing

Strategy: Send to all nodes (not just neighbors) information about directly connected links (not entire routing table).

- Link State Packet (LSP)
  - id of the node that created the LSP
  - cost of link to each directly connected neighbor
  - sequence number (SEQNO)
  - time-to-live (TTL) for this packet
- Reliable Flooding
  - store most recent LSP from each node
  - forward LSP to all nodes but one that sent it
  - generate new LSP periodically; increment SEQNO
  - start SEQNO at 0 when reboot
  - decrement TTL of each stored LSP; discard when TTL=0

# Link State

## Reliable Flooding



Flooding of link-state packets. (a) LSP arrives at node X; (b) X floods LSP to A and C; (c) A and C flood LSP to B (but not X); (d) flooding is complete

# Shortest Path Routing

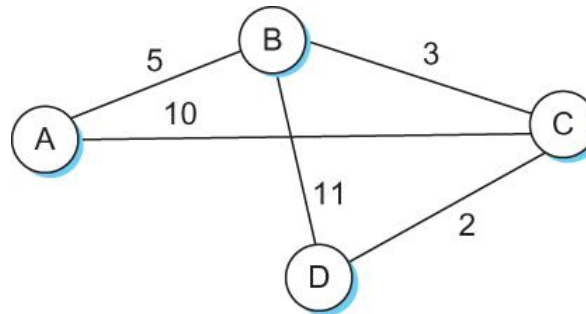
- In practice, each switch computes its routing table directly from the LSP's it has collected using a realization of Dijkstra's algorithm called the *forward search algorithm*
- Specifically each switch maintains two lists, known as **Tentative** and **Confirmed**
- Each of these lists contains a set of entries of the form (Destination, Cost, NextHop)

# Shortest Path Routing

## ■ The algorithm

- Initialize the **Confirmed** list with an entry for myself; this entry has a cost of 0
- For the node just added to the **Confirmed** list in the previous step, call it node **Next**, select its LSP
- For each neighbor (Neighbor) of **Next**, calculate the cost (Cost) to reach this Neighbor as the sum of the cost from myself to Next and from Next to Neighbor
  - If Neighbor is currently on neither the **Confirmed** nor the **Tentative** list, then add (Neighbor, Cost, Nexthop) to the **Tentative** list, where Nexthop is the direction I go to reach Next
  - If Neighbor is currently on the **Tentative** list, and the Cost is less than the currently listed cost for the Neighbor, then replace the current entry with (Neighbor, Cost, Nexthop) where Nexthop is the direction I go to reach Next
- If the **Tentative** list is empty, stop. Otherwise, pick the entry from the **Tentative** list with the lowest cost, move it to the **Confirmed** list, and return to Step 2.

# Shortest Path Routing



Step	Confirmed	Tentative	Comments
1	(D,0,-)		Since D is the only new member of the confirmed list, look at its LSP.
2	(D,0,-)	(B,11,B) (C,2,C)	D's LSP says we can reach B through B at cost 11, which is better than anything else on either list, so put it on Tentative list; same for C.
3	(D,0,-) (C,2,C)	(B,11,B)	Put lowest-cost member of Tentative (C) onto Confirmed list. Next, examine LSP of newly confirmed member (C).
4	(D,0,-) (C,2,C)	(B,5,C) (A,12,C)	Cost to reach B through C is 5, so replace (B,11,B). C's LSP tells us that we can reach A at cost 12.
5	(D,0,-) (C,2,C) (B,5,C)	(A,12,C)	Move lowest-cost member of Tentative (B) to Confirmed, then look at its LSP.
6	(D,0,-) (C,2,C) (B,5,C)	(A,10,C)	Since we can reach A at cost 5 through B, replace the Tentative entry.
7	(D,0,-) (C,2,C) (B,5,C) (A,10,C)		Move lowest-cost member of Tentative (A) to Confirmed, and we are all done.



# Shortest Path Routing

- Dijkstra's Algorithm - Assume non-negative link weights
  - $N$ : set of nodes in the graph
  - $l((i, j))$ : the non-negative cost associated with the edge between nodes  $i, j \in N$  and  $l(i, j) = \infty$  if no edge connects  $i$  and  $j$
  - Let  $s \in N$  be the starting node which executes the algorithm to find shortest paths to all other nodes in  $N$
  - Two variables used by the algorithm
    - $M$ : set of nodes incorporated so far by the algorithm
    - $C(n)$ : the cost of the path from  $s$  to each node  $n$
    - The algorithm

$M = \{s\}$

For each  $n$  in  $N - \{s\}$

$C(n) = l(s, n)$

while ( $N \neq M$ )

$M = M \cup \{w\}$  such that  $C(w)$  is the minimum  
for all  $w$  in  $(N-M)$

For each  $n$  in  $(N-M)$

$C(n) = \text{MIN} (C(n), C(w) + l(w, n))$

# **EXERCISE 63, PAGE 300**

# Today's Class

# IPv6

# Major Features

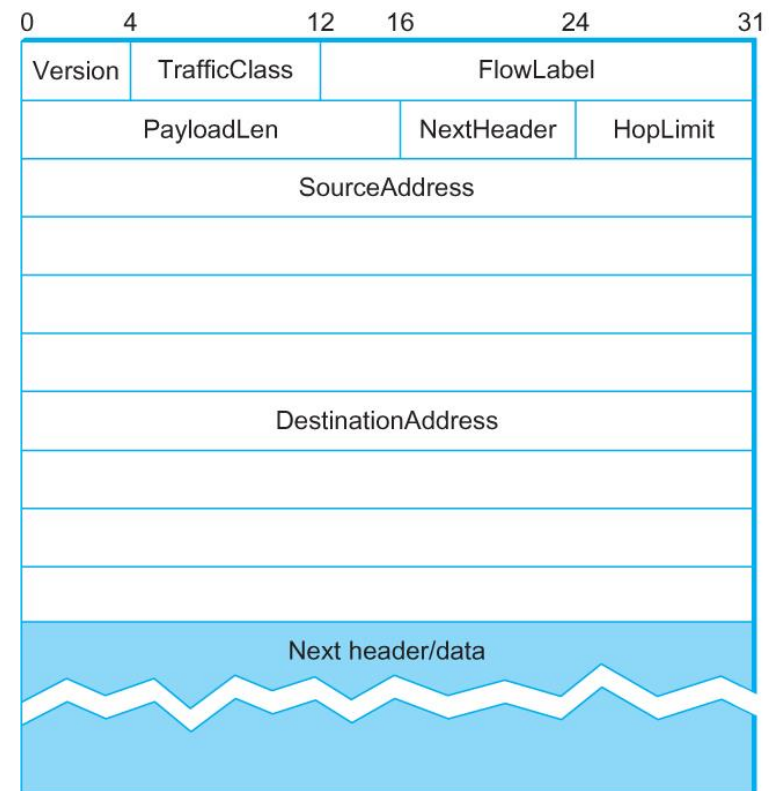
- 128-bit addresses
- Multicast
- Real-time service
- Authentication and security
- Auto-configuration
- End-to-end fragmentation
- Enhanced routing functionality, including support for mobile hosts

# IPv6 Addresses

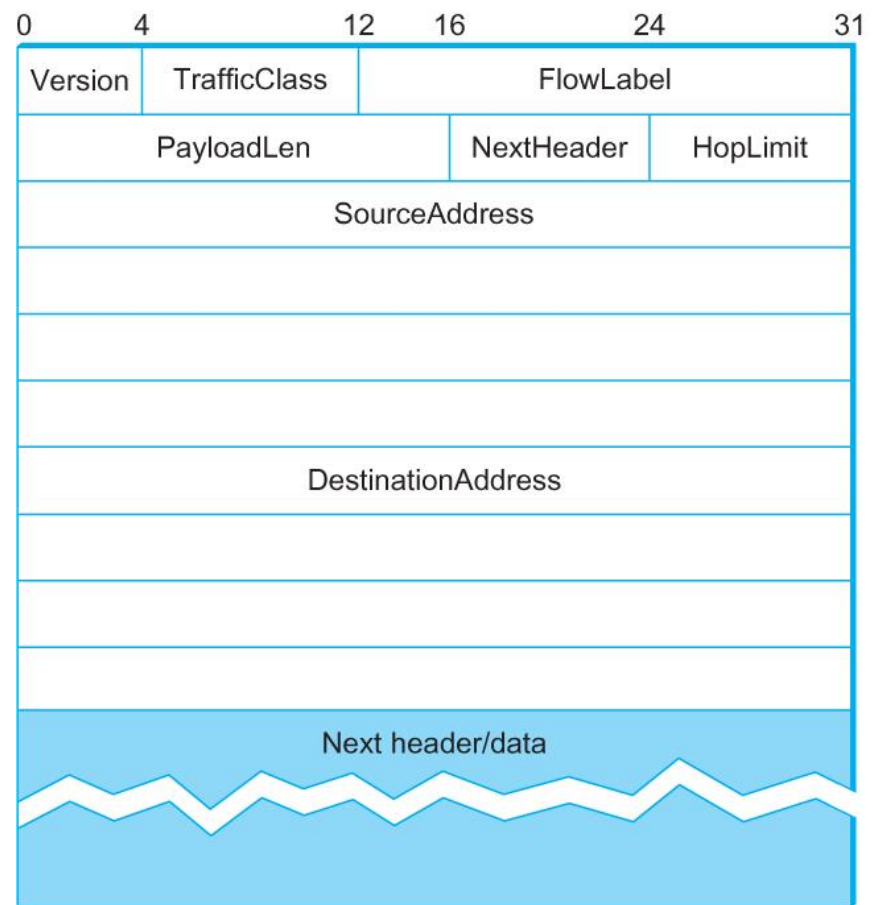
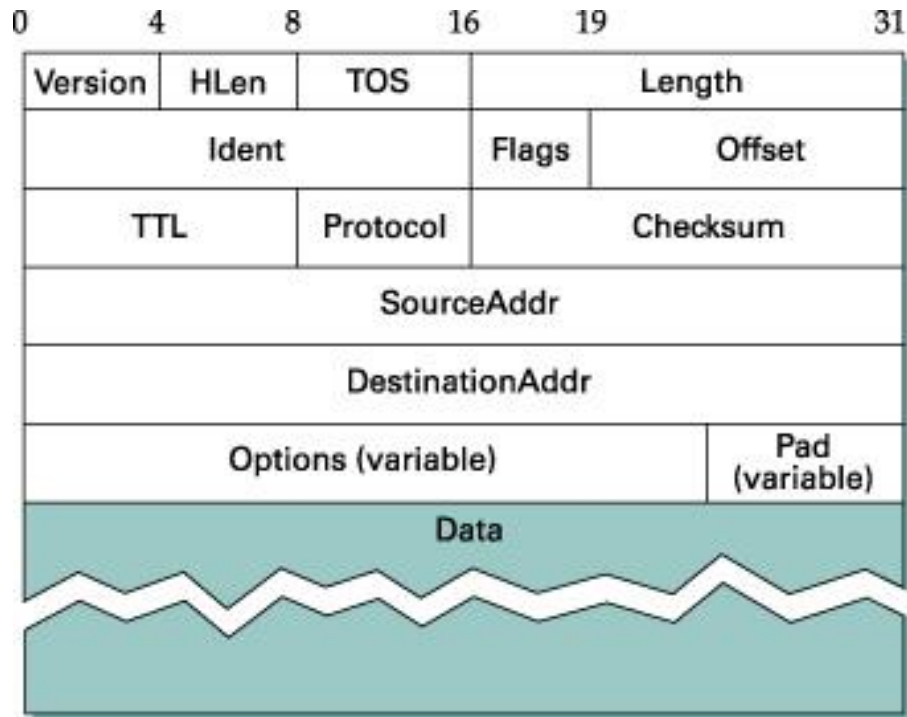
- Classless addressing/routing (similar to CIDR)
- Notation: x:x:x:x:x:x:x:x (x = 16-bit hex number)
  - contiguous 0s are compressed:  
47CD::A456:0124
  - IPv6 compatible IPv4 address: ::128.42.1.87
- Address assignment
  - provider-based
  - geographic

# IPv6 Header

- 40-byte “base” header
- Extension headers (fixed order, mostly fixed length)
  - fragmentation
  - source routing
  - authentication and security
  - other options



# IPv4 versus IPv6





# End-to-end Protocols

- Common properties that a transport protocol can be expected to provide
  - Guarantees message delivery
  - Delivers messages in the same order they were sent
  - Delivers at most one copy of each message
  - Supports arbitrarily large messages
  - Supports synchronization between the sender and the receiver
  - Allows the receiver to apply flow control to the sender
  - Supports multiple application processes on each host

# End-to-end Protocols

- Typical limitations of the network on which transport protocol will operate
  - Drop messages
  - Reorder messages
  - Deliver duplicate copies of a given message
  - Limit messages to some finite size
  - Deliver messages after an arbitrarily long delay

# End-to-end Protocols

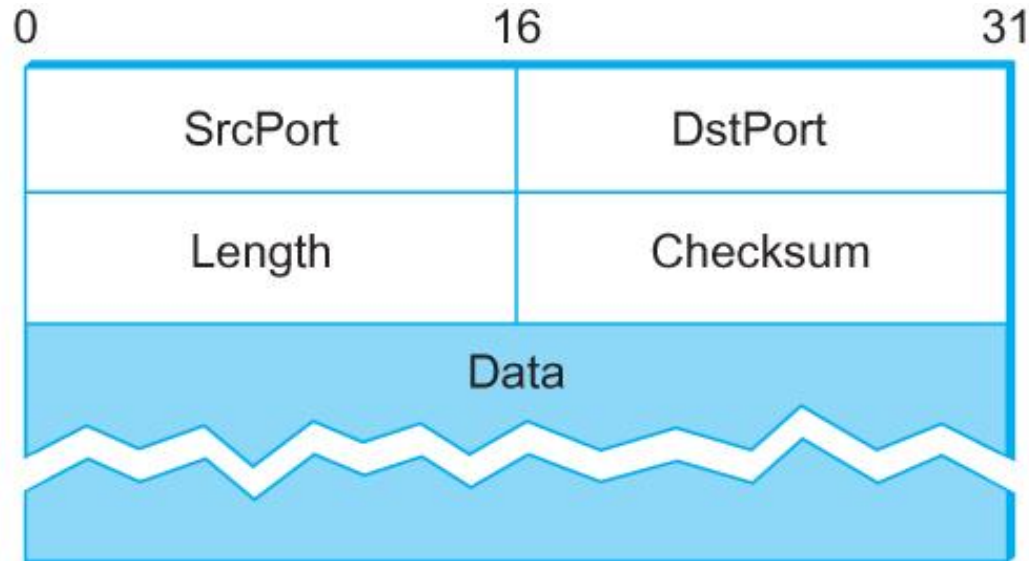
# End-to-end Protocols

- Challenge for Transport Protocols
  - Develop algorithms that turn the less-than-desirable properties of the underlying network into the high level of service required by application programs

# Simple Demultiplexer (UDP)

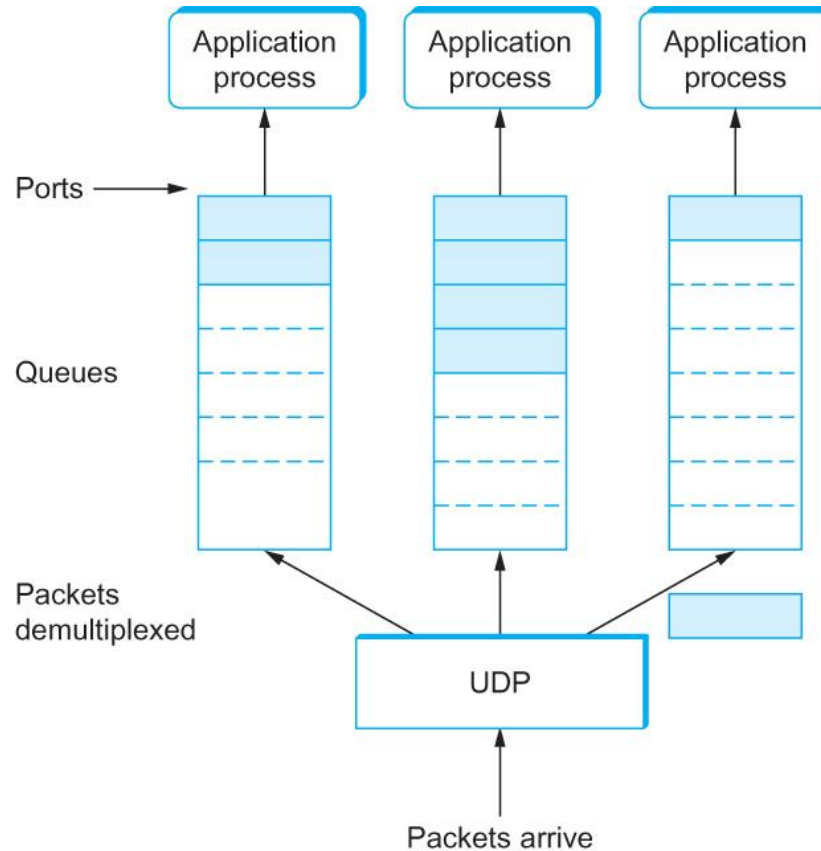
- Extends host-to-host delivery service of the underlying network into a process-to-process communication service
- Adds a level of demultiplexing which allows multiple application processes on each host to share the network

# Simple Demultiplexer (UDP)



Format for UDP header (Note: length and checksum fields should be switched)

# Simple Demultiplexer (UDP)



UDP Message Queue

# Reliable Byte Stream (TCP)

- In contrast to UDP, Transmission Control Protocol (TCP) offers the following services
  - Reliable
  - Connection oriented
  - Byte-stream service



# Flow control VS Congestion control

- Flow control involves preventing senders from overrunning the capacity of the receivers
- Congestion control involves preventing too much data from being injected into the network, thereby causing switches or links to become overloaded

# End-to-end Issues

- At the heart of TCP is the sliding window algorithm (discussed in Chapter 2)
- As TCP runs over the Internet rather than a point-to-point link, the following issues need to be addressed by the sliding window algorithm
  - TCP supports logical connections between processes that are running on two different computers in the Internet
  - TCP connections are likely to have widely different RTT times
  - Packets may get reordered in the Internet

# End-to-end Issues

- TCP needs a mechanism using which each side of a connection will learn what resources the other side is able to apply to the connection
- TCP needs a mechanism using which the sending side will learn the capacity of the network

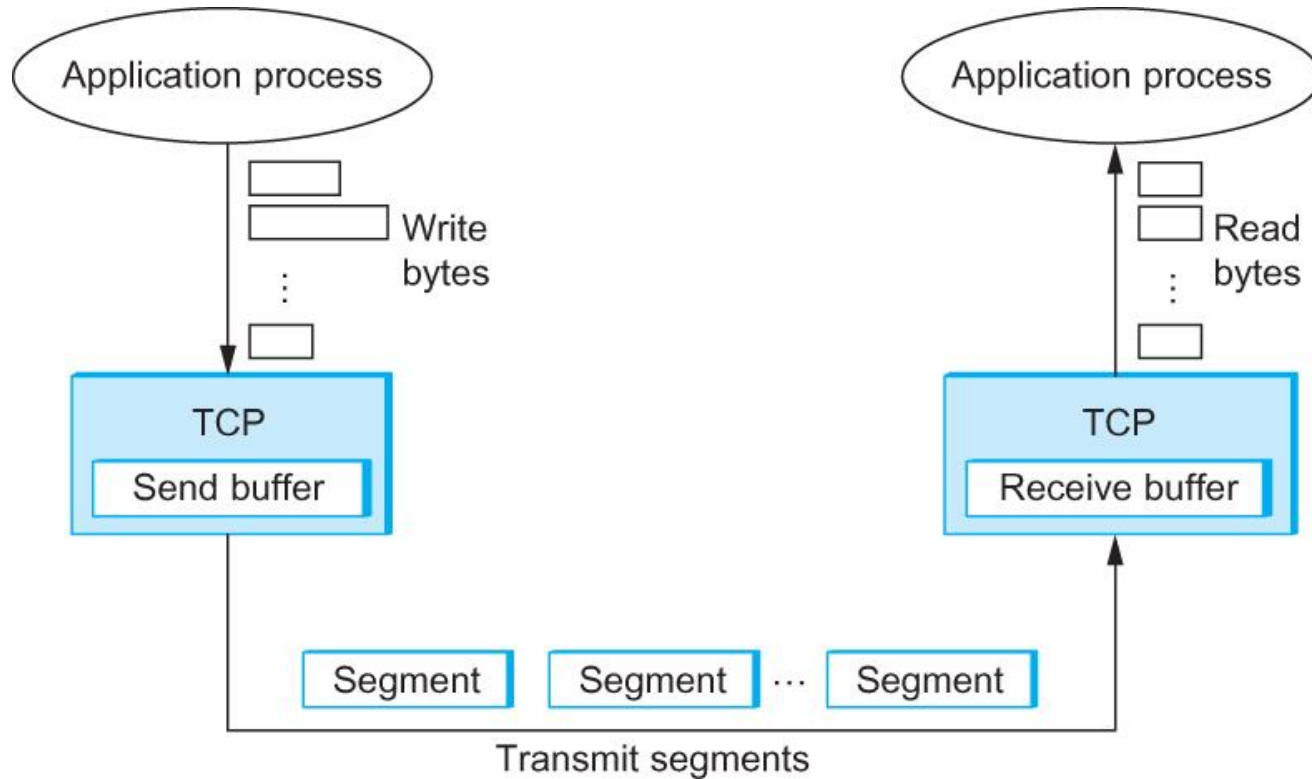
# TCP Segment

- TCP is a byte-oriented protocol, which means that the sender writes bytes into a TCP connection and the receiver reads bytes out of the TCP connection.
- Although “byte stream” describes the service TCP offers to application processes, TCP does not, itself, transmit individual bytes over the Internet.

# TCP Segment

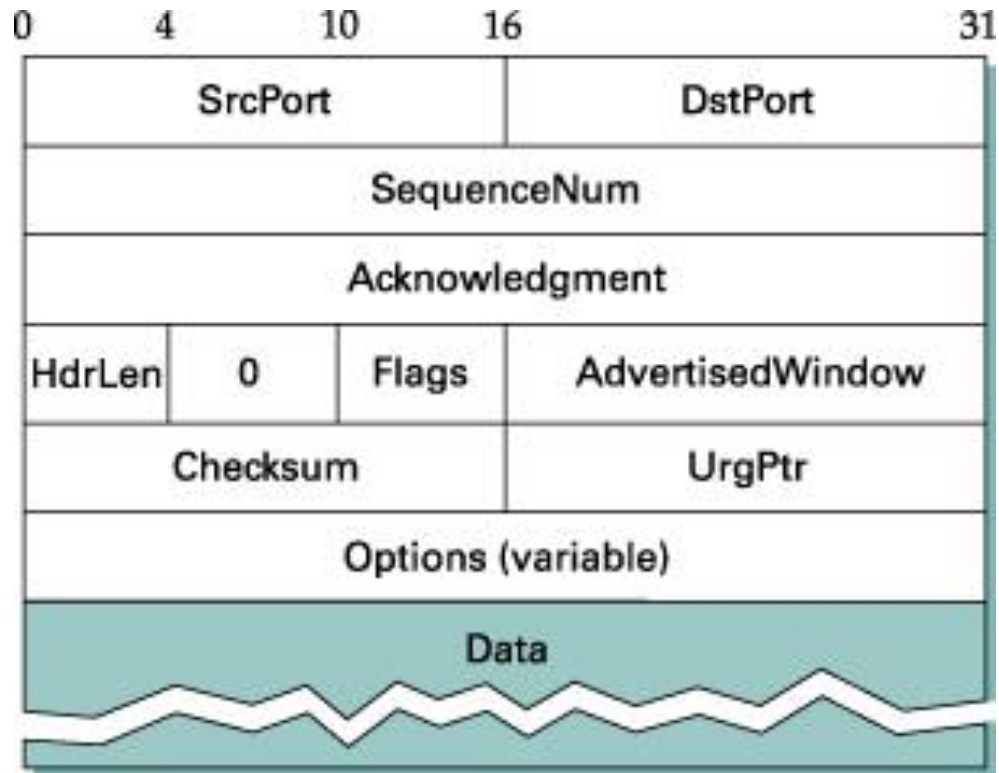
- TCP on the source host buffers enough bytes from the sending process to fill a reasonably sized packet and then sends this packet to its peer on the destination host.
- TCP on the destination host then empties the contents of the packet into a receive buffer, and the receiving process reads from this buffer at its leisure.
- The packets exchanged between TCP peers are called *segments*.

# TCP Segment

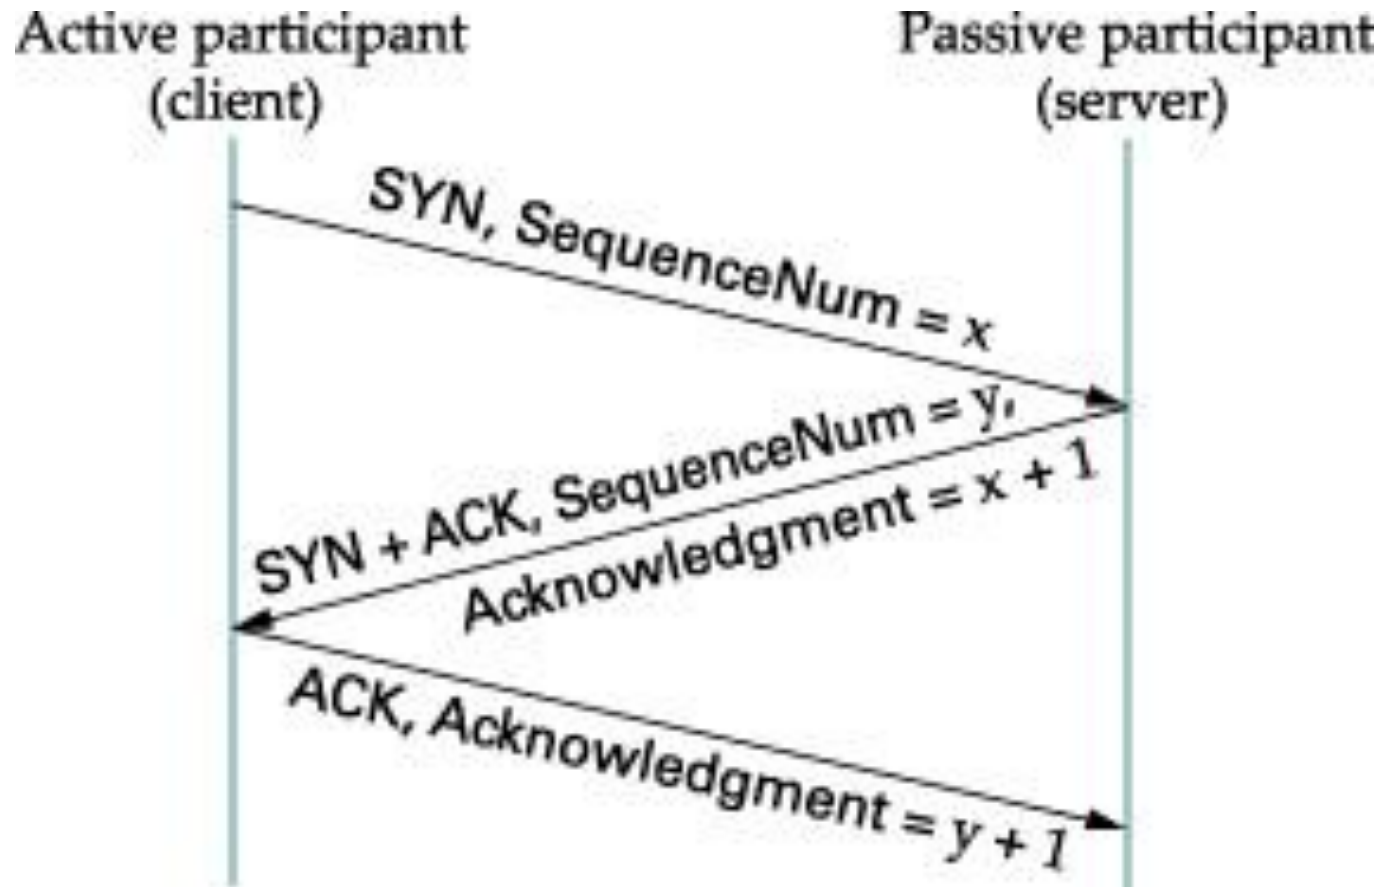


How TCP manages a byte stream.

# TCP Packet



# Connection Establishment in TCP





# Connection Establishment in TCP

sample\_data01.tcpdump.gz [Wireshark 1.6.2 (SVN Rev 38931 from /trunk-1.6)]

File Edit View Go Capture Analyze Statistics Telephony Tools Internals Help

Filter: tcp Expression... Clear Apply

No.	Time	Source	Destination	Protocol	Length	Info
13527	529.300055	192.168.0.40	192.168.1.30	TELNET	60	Telnet Data ...
13528	529.362511	192.168.1.30	192.168.0.40	TELNET	60	Telnet Data ...
13529	529.410377	192.168.0.40	192.168.1.30	TCP	60	43612 > telnet [ACK] Seq=368 Ack=2184 Win=8760 Len=0
13530	529.500696	192.168.0.40	192.168.1.30	TELNET	60	Telnet Data ...
13531	529.502397	192.168.1.30	192.168.0.40	TELNET	60	Telnet Data ...
13532	529.550421	192.168.0.40	192.168.1.30	TCP	60	43612 > telnet [ACK] Seq=369 Ack=2185 Win=8760 Len=0
13533	529.625458	192.168.1.30	192.168.0.20	TCP	60	afrog > telnet [SYN] Seq=0 Win=512 Len=0 MSS=1460
13534	529.627681	192.168.0.20	192.168.1.30	TCP	60	telnet > afrog [SYN, ACK] Seq=0 Ack=1 Win=8760 Len=0 MSS=1460
13535	529.628121	192.168.1.30	192.168.0.20	TCP	60	afrog > telnet [ACK] Seq=1 Ack=1 Win=31744 Len=0
13536	529.640712	192.168.0.40	192.168.1.30	TELNET	60	Telnet Data ...
13537	529.642037	192.168.1.30	192.168.0.40	TELNET	60	Telnet Data ...
13538	529.646300	192.168.1.30	192.168.0.20	TELNET	84	Telnet Data ...

Frame 13538: 84 bytes on wire (672 bits), 84 bytes captured (672 bits)

Ethernet II, Src: 3com\_79:af:be (00:60:97:79:af:be), Dst: Cisco\_04:41:bc (00:00:0c:04:41:bc)

Internet Protocol Version 4, Src: 192.168.1.30 (192.168.1.30), Dst: 192.168.0.20 (192.168.0.20)

Transmission Control Protocol, Src Port: afrog (1042), Dst Port: telnet (23), Seq: 1, Ack: 1, Len: 30

Telnet

0000 00 00 0c 04 41 bc 00 60 97 79 af be 08 00 45 00 . . . . A . . . y . . . . E .  
0010 00 46 db 95 40 00 40 06 dc 99 c0 a8 01 1e c0 a8 . F . . @ . @ . . . . .  
0020 00 14 04 12 00 17 d0 a3 49 a9 2f e0 86 ed 50 18 . . . . . I . / . . . P .  
0030 7c 00 69 6c 00 00 ff fb 25 ff fd 03 ff fb 18 ff | . i l . . . . % . . . . .  
0040 fb 1f ff fb 20 ff fb 21 ff fb 22 ff fb 27 ff fd . . . . . ! . . . . . ' . . . .

File: "C:\Fall 2011\CSCL 460\TCP Du... Packets: 14523 Displayed: 14282 Marked: 2 Load time: 0:00.233 Profile: Default

8:54 PM 9/26/2011

# Connection Tear-down in TCP

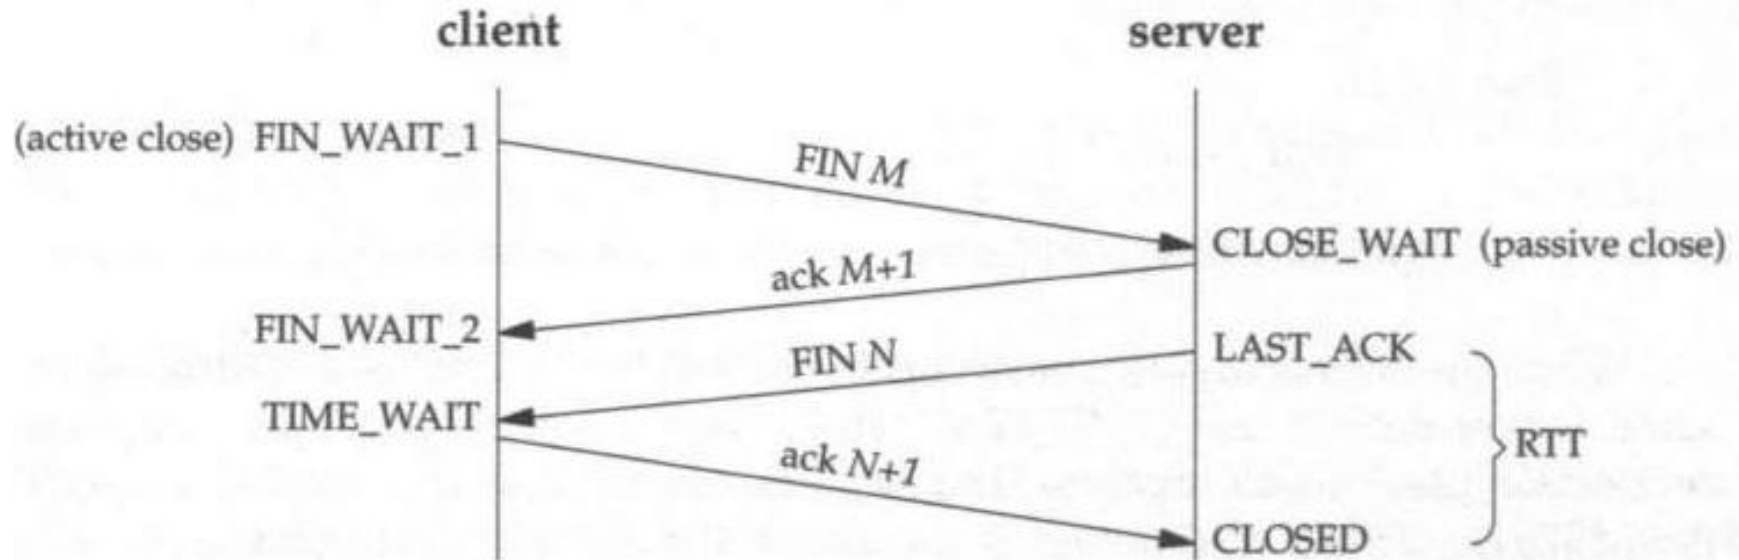
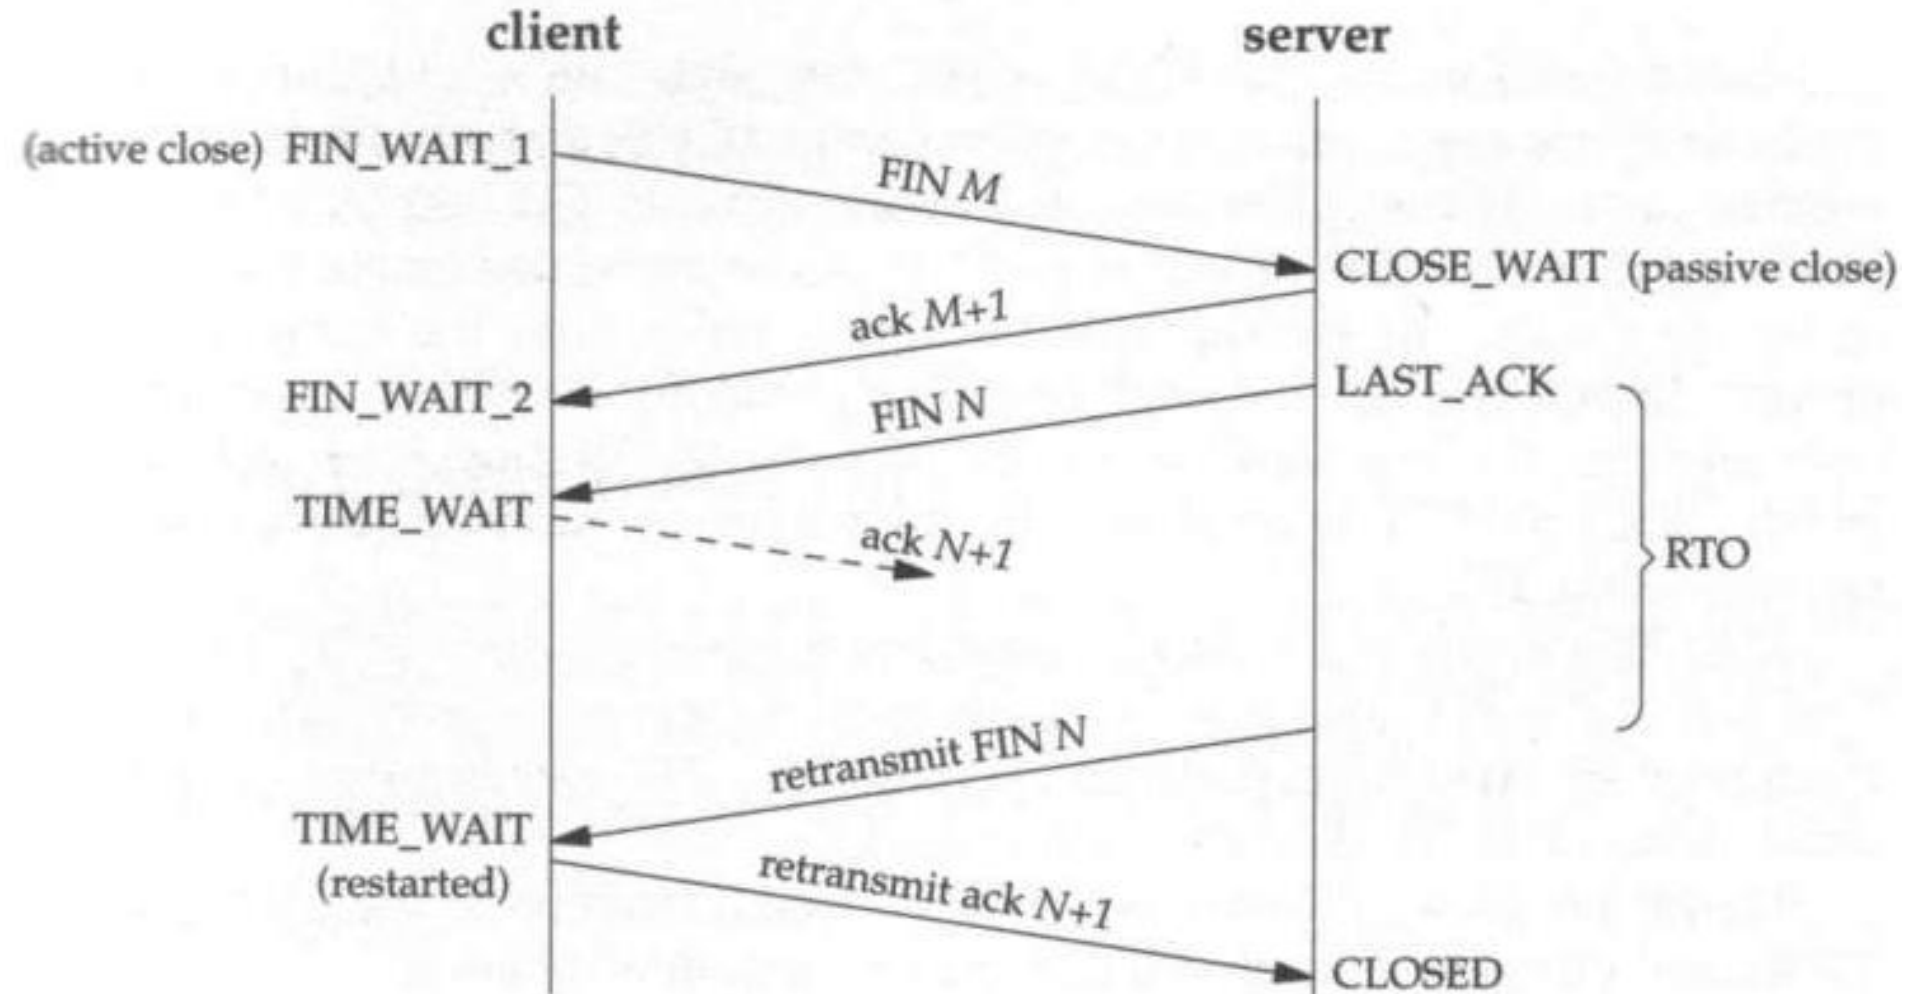
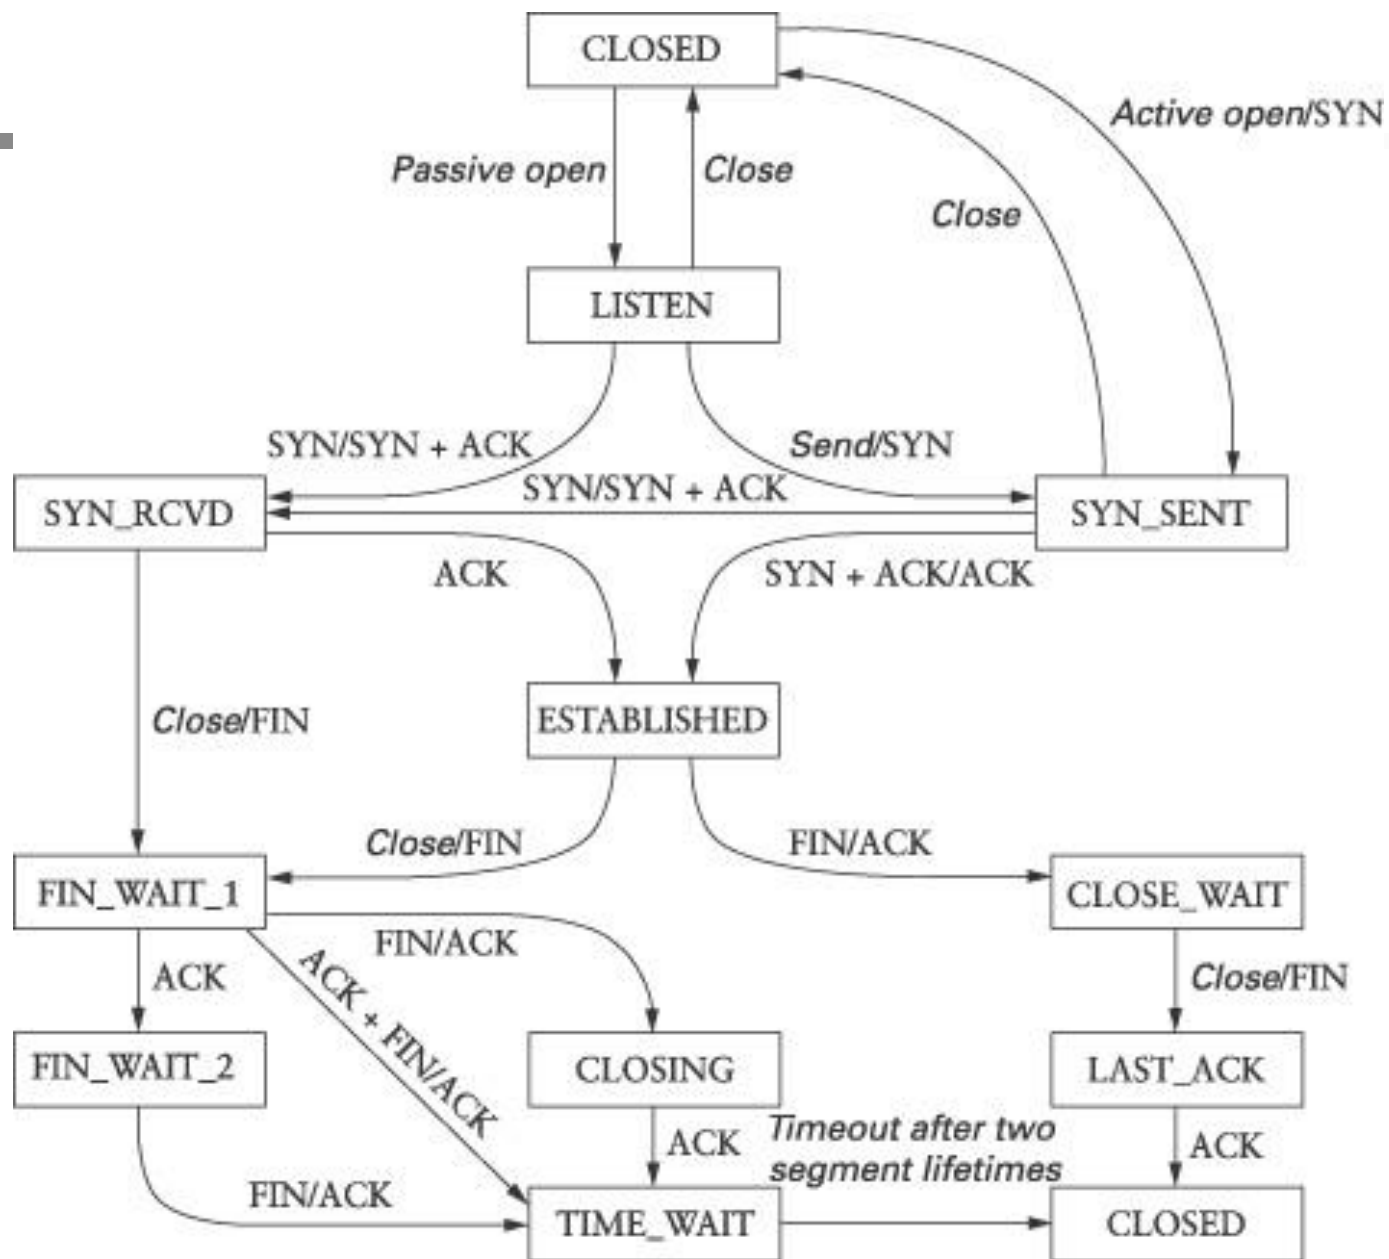


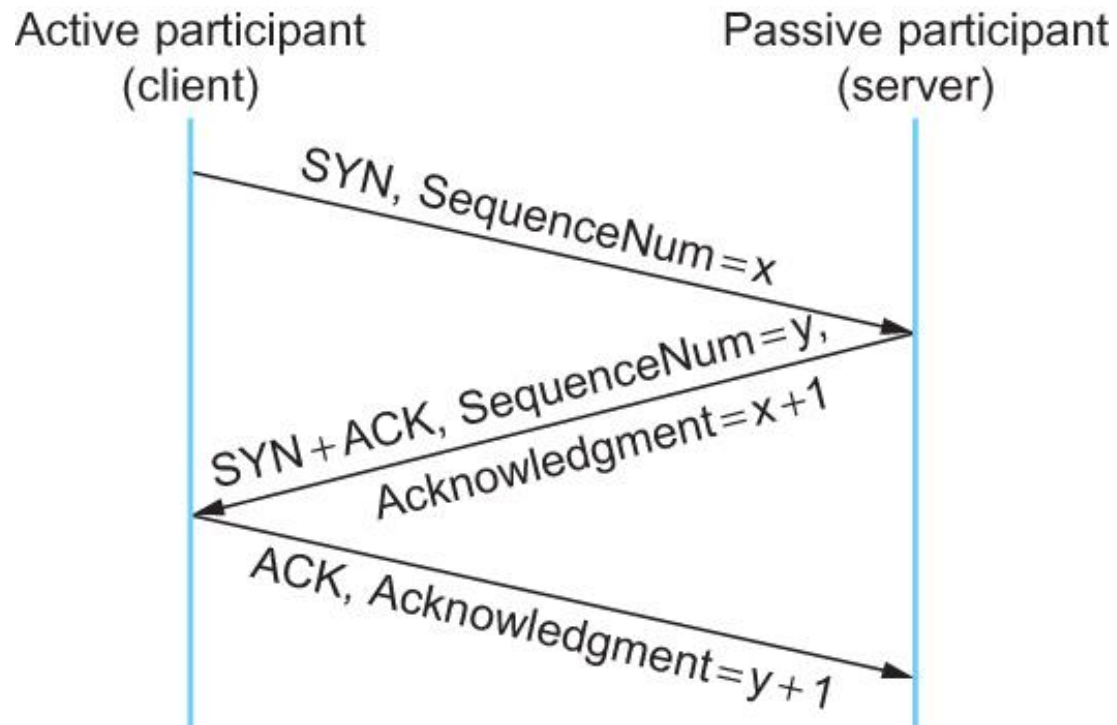
Figure 4.4 Normal exchange of segments to close a connection.

# Connection Tear-down in TCP





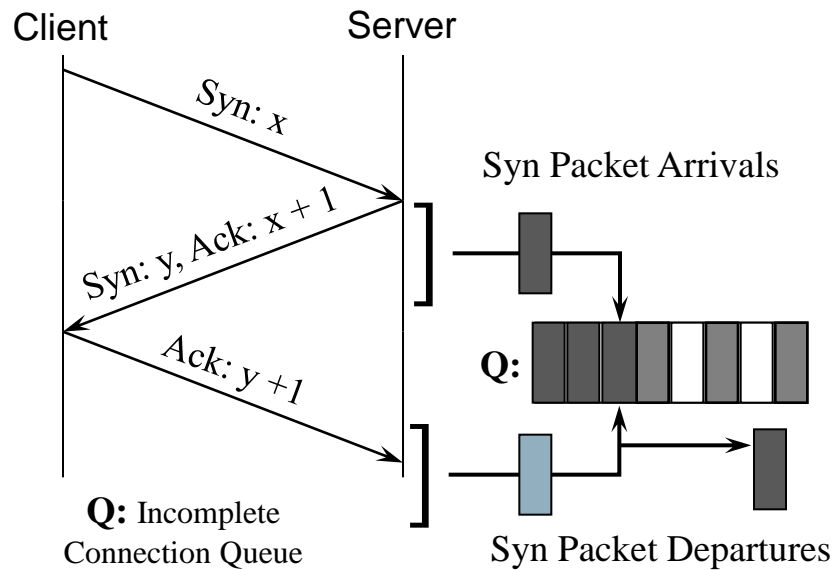
# Connection Establishment in TCP



Timeline for three-way handshake algorithm

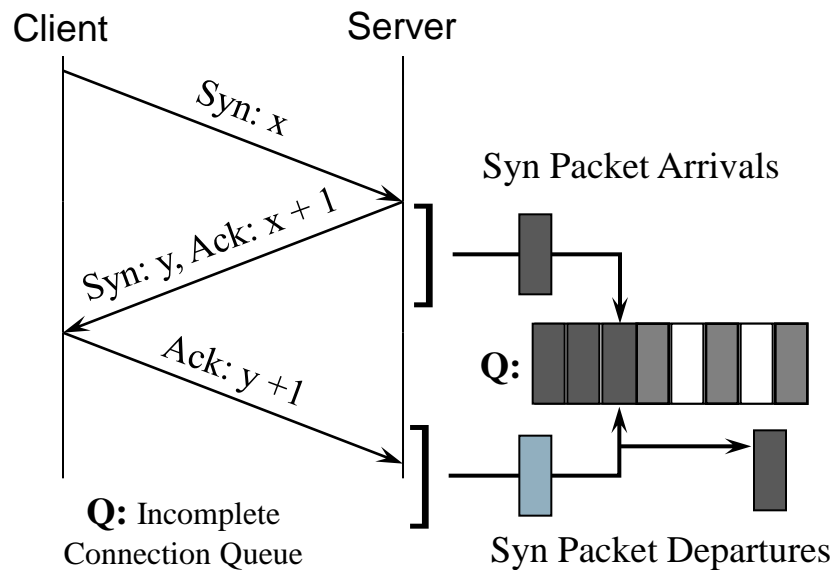
# SYN Flooding in TCP

## Normal Three-way Handshake in TCP

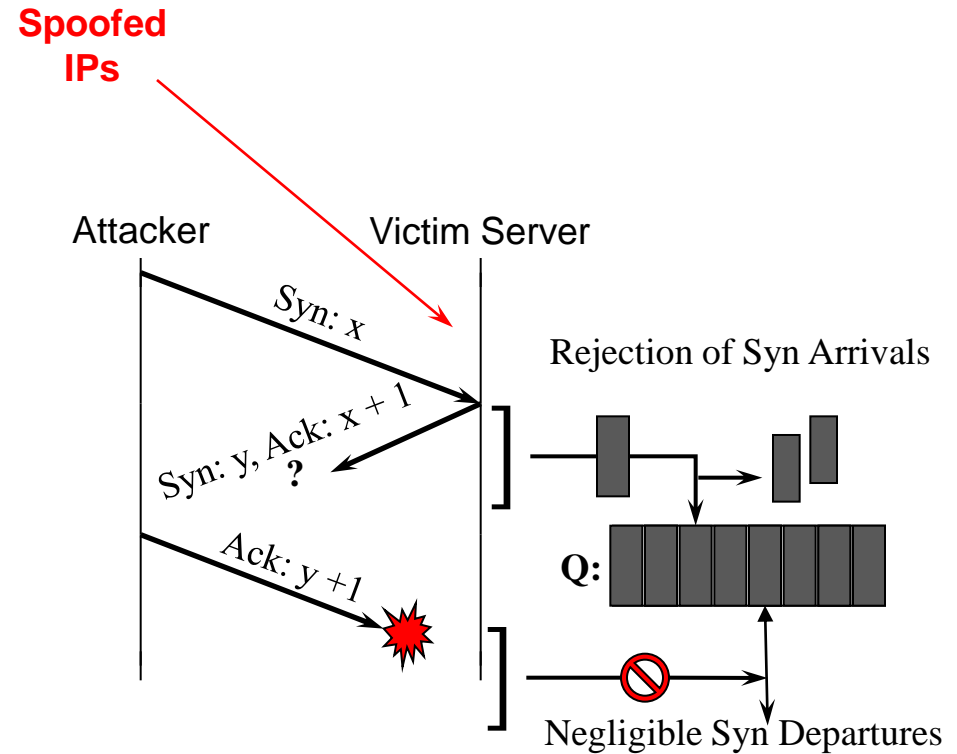


# SYN Flooding in TCP

## Normal Three-way Handshake in TCP



## SYN Flooding Attack



**Attack Signature**

# SYN Flooding in TCP

tcpdump.gz [Wireshark 1.6.2 (SVN Rev 38931 from /trunk-1.6)]

File Edit View Go Capture Analyze Statistics Telephony Tools Internals Help

Filter: tcp Expression... Clear Apply

No.	Time	Source	Destination	Protocol	Length	Info
276729	11:55:15.067167	135.8.60.182	172.16.112.207	TCP	60	15045 > smtp [FIN, ACK] Seq=997 Ack=329 Win=32120 Len=0
276730	11:55:15.067464	135.8.60.182	172.16.112.194	TCP	60	15051 > smtp [SYN] Seq=0 Win=512 Len=0 MSS=1460
276731	11:55:15.067588	172.16.112.207	135.8.60.182	TCP	60	smtp > 15045 [ACK] Seq=329 Ack=998 Win=32735 Len=0
276732	11:55:15.067967	172.16.112.194	135.8.60.182	TCP	60	smtp > 15051 [SYN, ACK] Seq=0 Ack=1 Win=32736 Len=0 MSS=1460
276733	11:55:15.068120	135.8.60.182	172.16.112.194	TCP	60	15051 > smtp [ACK] Seq=1 Ack=1 Win=32120 Len=0
276734	11:55:15.453647	1.2.3.4	172.16.112.50	TCP	60	18700 > telnet [SYN] Seq=0 Win=242 Len=0
276735	11:55:15.473565	1.2.3.4	172.16.112.50	TCP	60	18956 > telnet [SYN] Seq=0 Win=242 Len=0
276736	11:55:15.493560	1.2.3.4	172.16.112.50	TCP	60	19212 > telnet [SYN] Seq=0 Win=242 Len=0
276737	11:55:15.513563	1.2.3.4	172.16.112.50	TCP	60	19468 > telnet [SYN] Seq=0 Win=242 Len=0
276738	11:55:15.533564	1.2.3.4	172.16.112.50	TCP	60	19724 > telnet [SYN] Seq=0 Win=242 Len=0
276739	11:55:15.553575	1.2.3.4	172.16.112.50	TCP	60	19980 > telnet [SYN] Seq=0 Win=242 Len=0
276740	11:55:15.573562	1.2.3.4	172.16.112.50	TCP	60	20236 > telnet [SYN] Seq=0 Win=242 Len=0
276741	11:55:15.593562	1.2.3.4	172.16.112.50	TCP	60	20492 > telnet [SYN] Seq=0 Win=242 Len=0
276742	11:55:15.613564	1.2.3.4	172.16.112.50	TCP	60	20748 > telnet [SYN] Seq=0 Win=242 Len=0
276743	11:55:15.633565	1.2.3.4	172.16.112.50	TCP	60	21004 > telnet [SYN] Seq=0 Win=242 Len=0
276744	11:55:15.653571	1.2.3.4	172.16.112.50	TCP	60	21260 > telnet [SYN] Seq=0 Win=242 Len=0
276745	11:55:15.673570	1.2.3.4	172.16.112.50	TCP	60	21516 > telnet [SYN] Seq=0 Win=242 Len=0
276746	11:55:15.693713	1.2.3.4	172.16.112.50	TCP	60	21772 > telnet [SYN] Seq=0 Win=242 Len=0
276749	11:55:15.786412	172.16.112.149	195.115.218.108	SMTP	137	S: 220 eagle.eyrie.af.mil ESMTP Sendmail 8.8.7/8.8.7; wed, 3 Jun 1998 11:55:15 -0400
276750	11:55:15.789426	195.115.218.108	172.16.112.149	SMTP	77	C: EHLO epsilon.pear.com
276751	11:55:15.799483	172.16.112.149	195.115.218.108	SMTP	80	S: 500 Command unrecognized
276752	11:55:15.799784	195.115.218.108	172.16.112.149	SMTP	77	C: HELO epsilon.pear.com
276753	11:55:15.800322	172.16.112.149	195.115.218.108	SMTP	99	S: 250 (epsilon.pear.com) pleased to meet you.
276754	11:55:15.800579	195.115.218.108	172.16.112.149	SMTP	91	C: MAIL From:<anguse@epsilon.pear.com>
276755	11:55:15.801080	172.16.112.149	195.115.218.108	SMTP	98	S: 250 <anguse@epsilon.pear.com>... Sender Ok
276756	11:55:15.801375	195.115.218.108	172.16.112.149	SMTP	93	C: RCPT To:<enriqueg@eagle.eyrie.af.mil>
276757	11:55:15.801868	172.16.112.149	195.115.218.108	SMTP	92	S: 250 <enriqueg@eagle.eyrie.af.mil> OK
276758	11:55:15.802130	195.115.218.108	172.16.112.149	SMTP	60	C: DATA
276759	11:55:15.802844	172.16.112.149	195.115.218.108	SMTP	104	C: 254 Enter mail, end with "." on a line by itself

Frame 276735: 60 bytes on wire (480 bits), 60 bytes captured (480 bits)

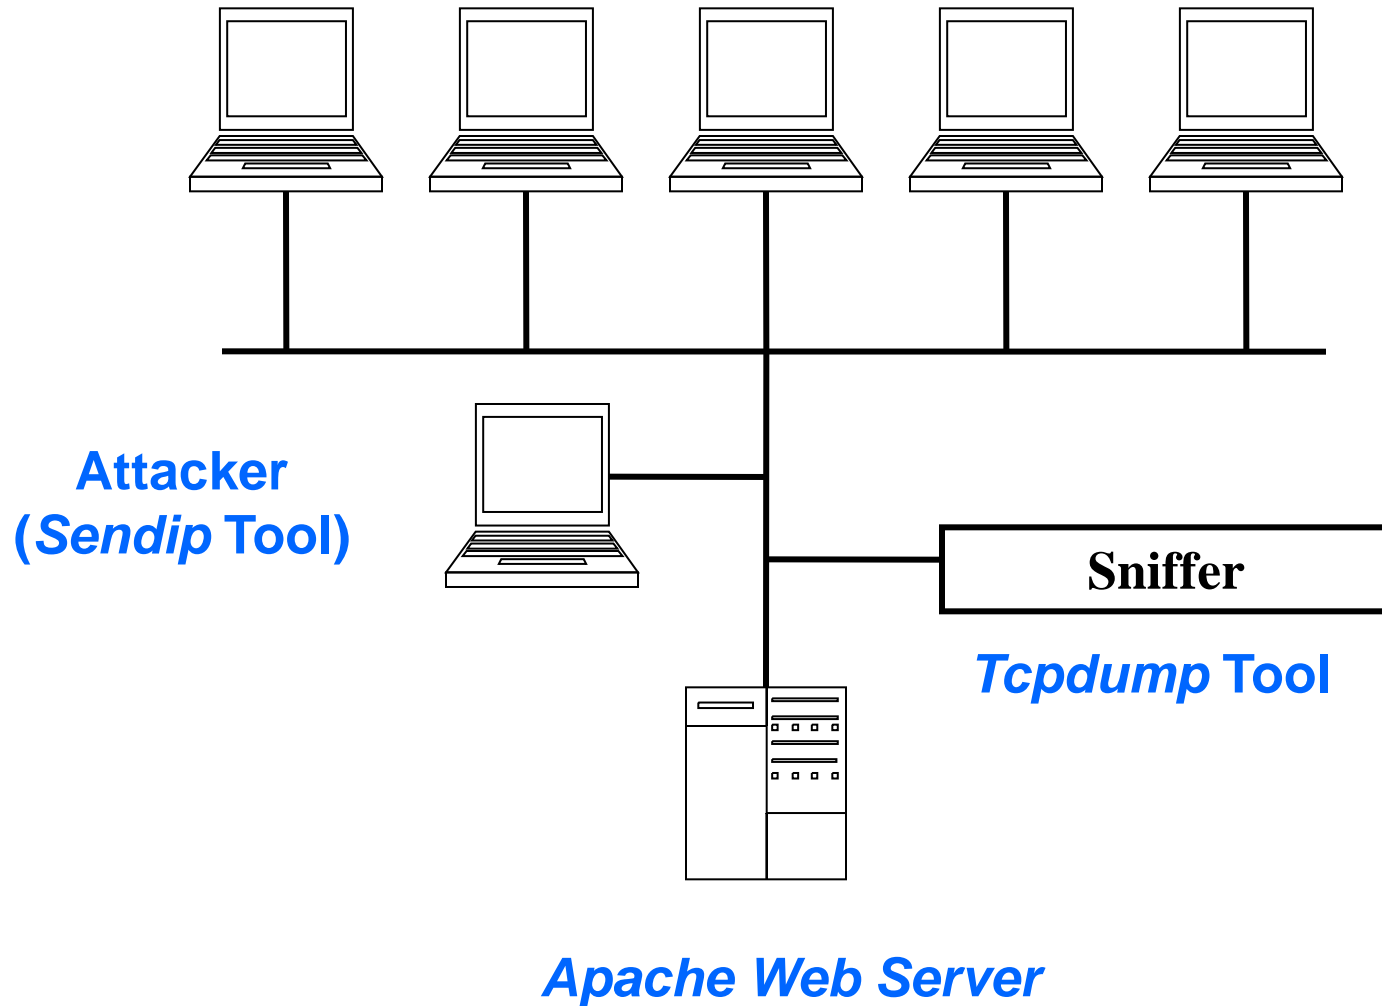
Ethernet II, Src: DellComp\_a3:58:23 (00:c0:4f:a3:58:23), Dst: Cisco\_04:41:bc (00:00:0c:04:41:bc)

Internet Protocol Version 4, Src: 1.2.3.4 (1.2.3.4), Dst: 172.16.112.50 (172.16.112.50)

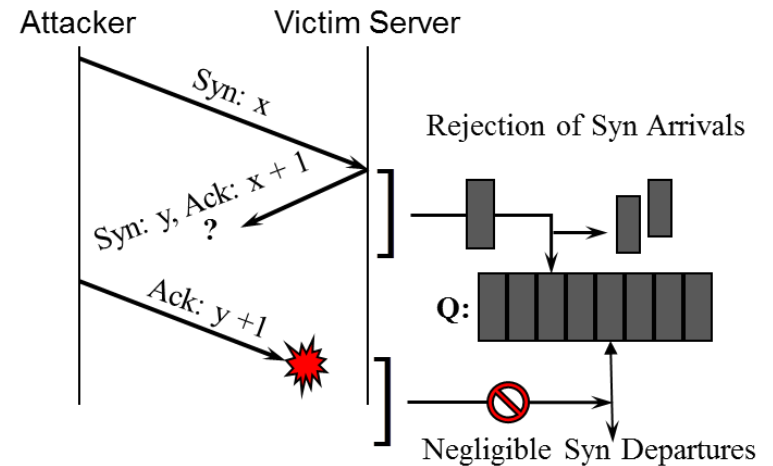
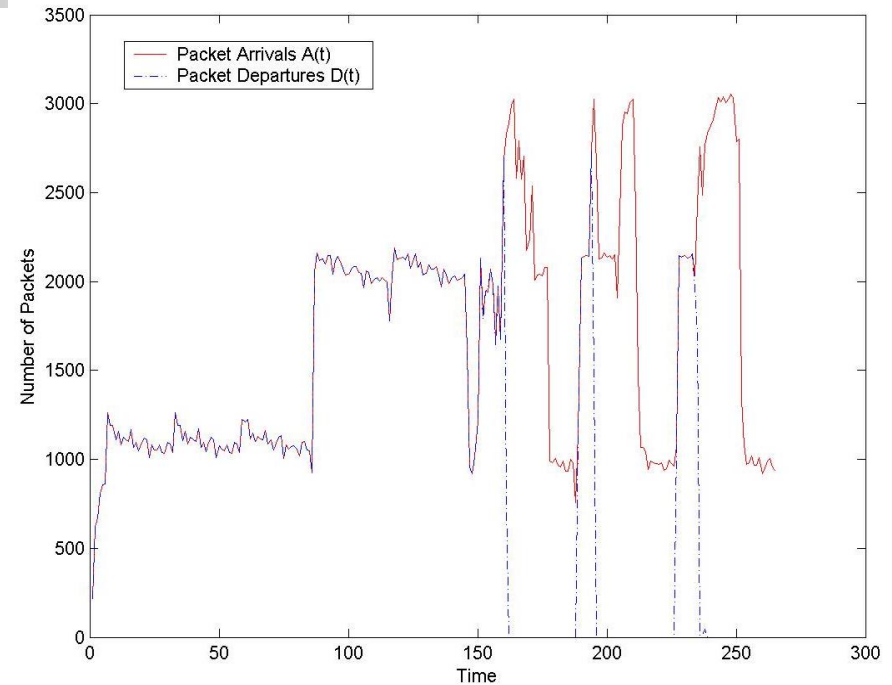
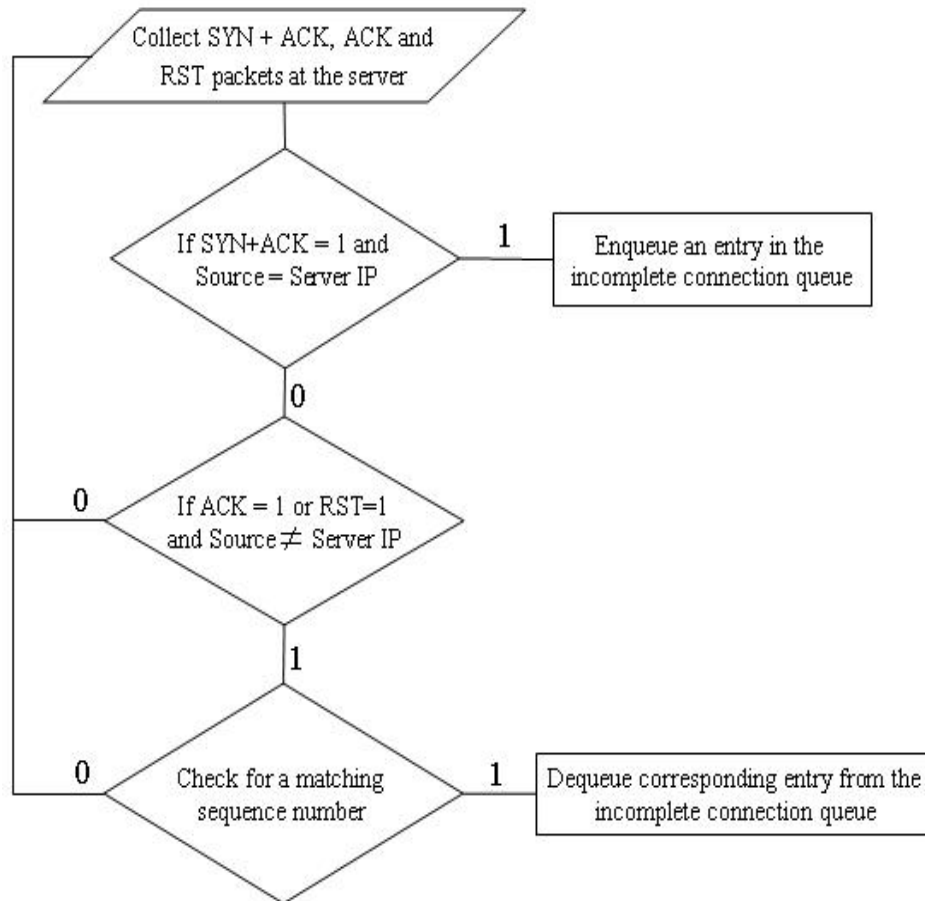
Transmission Control Protocol, Src Port: 18956 (18956), Dst Port: telnet (23), Seq: 0, Len: 0



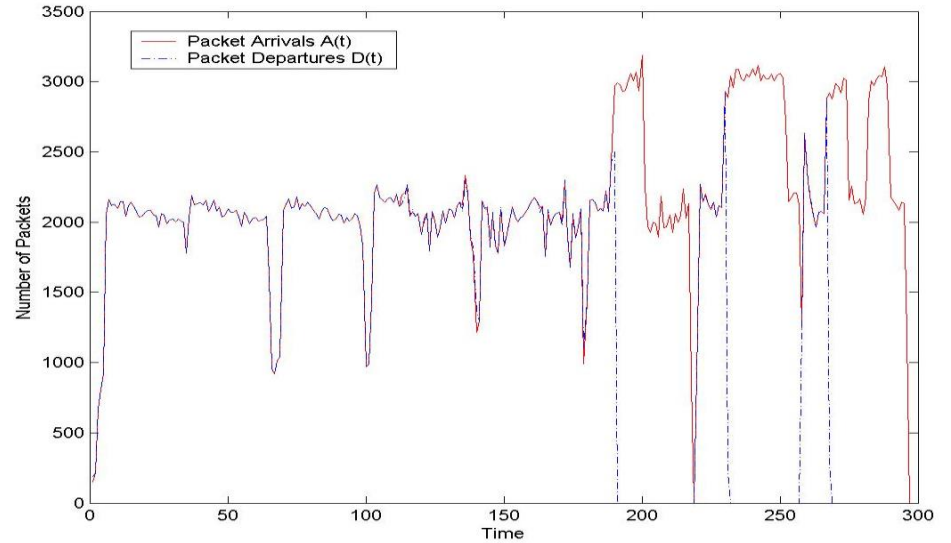
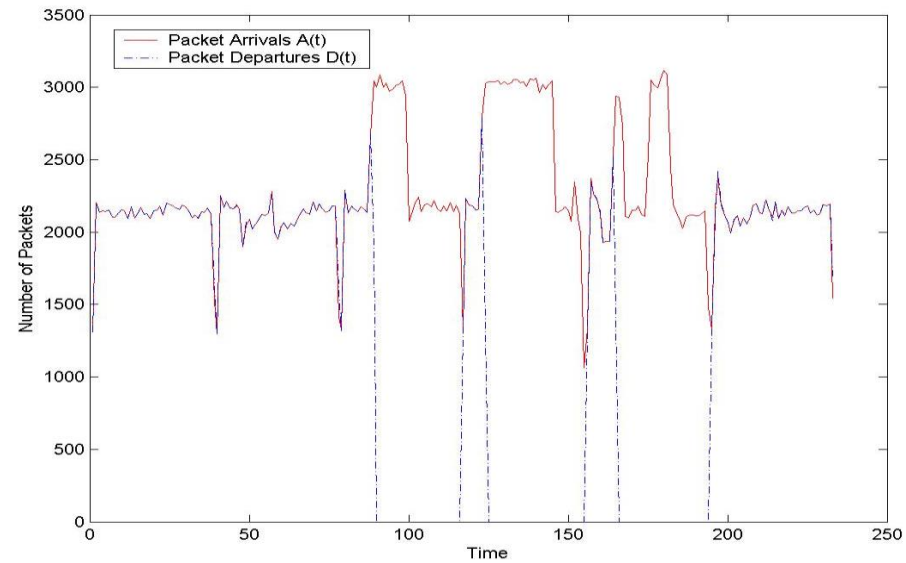
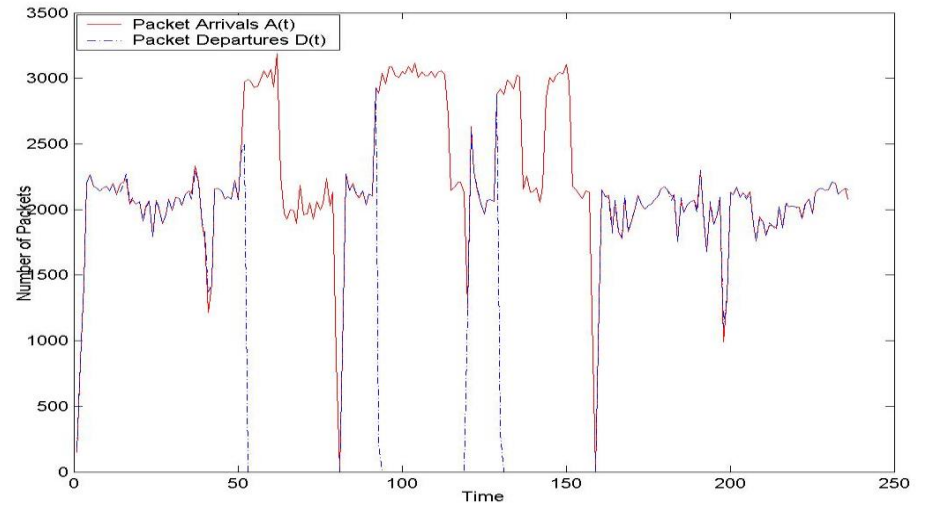
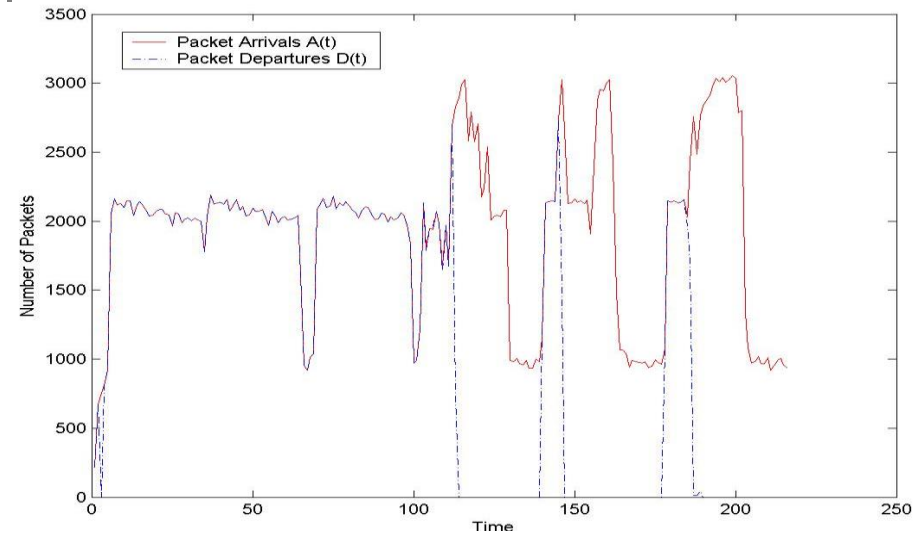
## Clients (*SURGE* Tool)



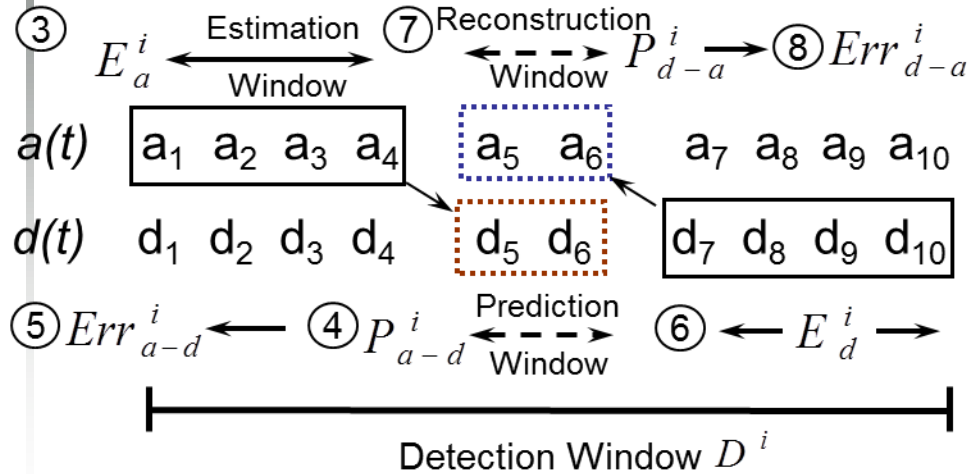
# Packet Classification and Counting



# Packet Counts



# SYNFloodAlert Algorithm



- Function Estimation by polynomial function approximation on  $P_3 = a + bx + cx^2 + dx^3$

- Error

$$Err^i = \max(f_a^i(t) - A(t)) \text{ or } \max(f_d^i(t) - D(t))$$

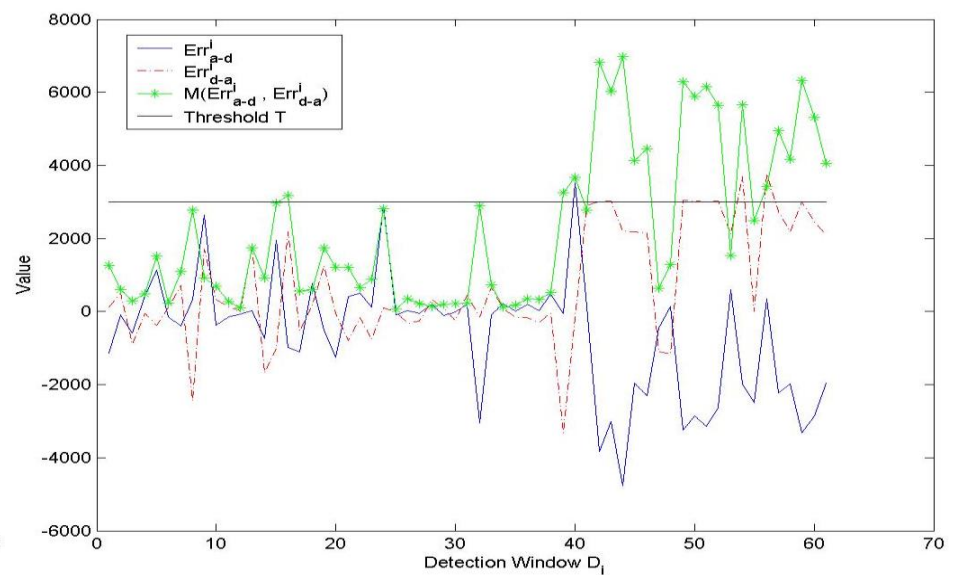
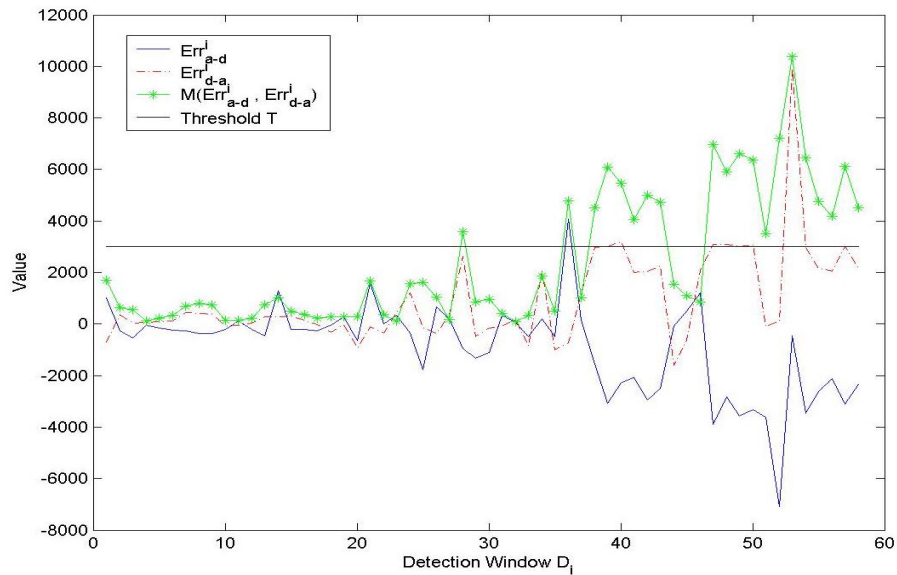
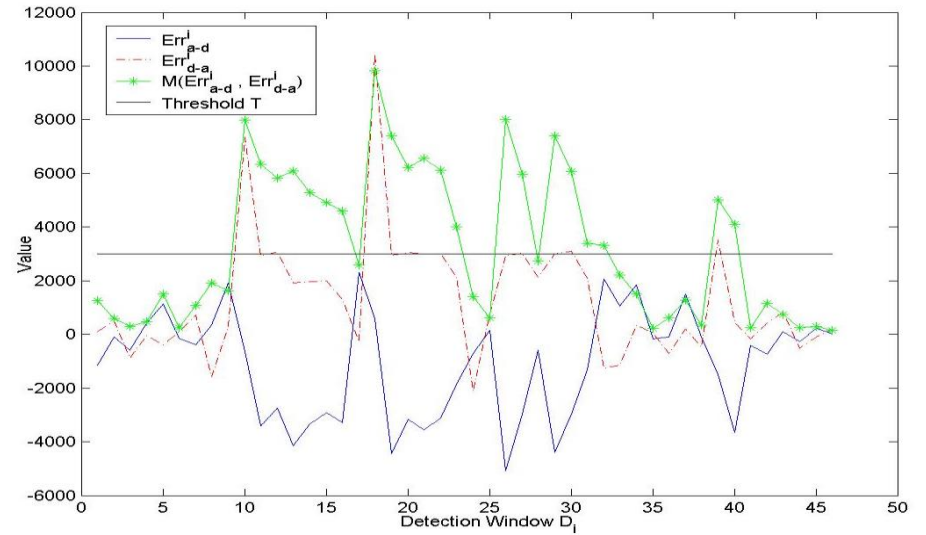
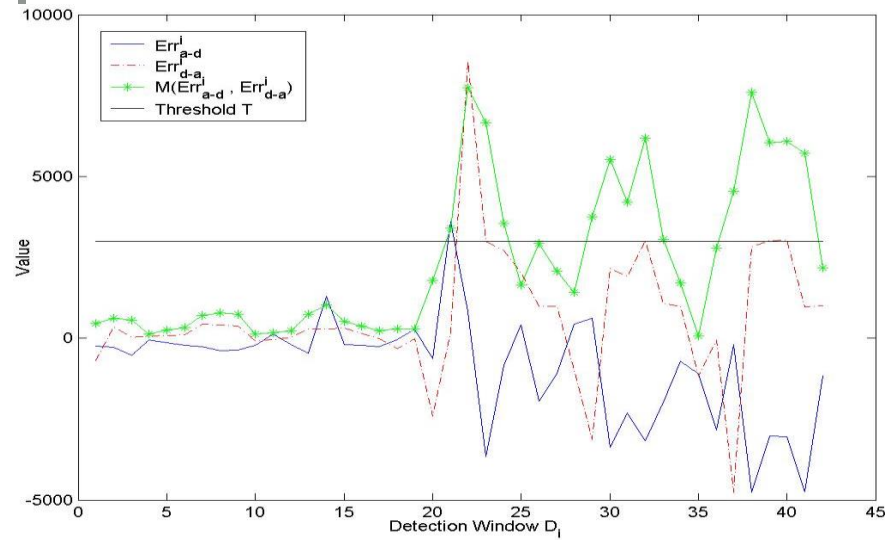
- Tunable Threshold  $T$

/\* Set  $Err_{a-d}^i$  and  $Err_{d-a}^i$  \*/

/\* Set the estimation, prediction, and reconstruction window sizes for Syn arrivals and departures. \*/

1. Estimate function  $f_a^i$  using Syn arrivals in the Arrival Estimation Window.
2. Predict Syn departures using  $f_a^i$  for the departures in the Prediction Window  $P_{a-d}^i$ . Compute error  $Err_{a-d}^i$ .
3. Estimate function  $f_d^i$  for Syn departures in the Departure Estimation Window.
4. Reconstruct previous SYN arrivals using  $f_d^i$  for all arrivals in the Reconstruction Window. Compute error  $P_{d-a}^i$ .
5. If  $D(Err_{a-d}^i - Err_{d-a}^i) > T$ , Flag an attack.

# SYNFloodAlert Algorithm



**THE END**

**$C \rightarrow S : SIN (ISN_c)$**

**$S \rightarrow C : SYN$**

**$(ISN_s), ACK (ISN_c)$**

**$C \rightarrow S : ACK (ISN_s)$**

**$C \rightarrow S : data$**

***and/or***

**$S \rightarrow C : data$**

**$X \rightarrow S : SYN (ISN_x), SRC = T$**

**$S \rightarrow T : SYN (ISN_s), ACK$   
 $(ISN_x)$**

**$X \rightarrow S : ACK (ISN_s), SRC = T$**

**$X \rightarrow S : ACK (ISN_s), SRC = T,$   
*nasty-data***



**Explain the sequence number prediction attack on TCP (see ‘Security Problems in the TCP/IP Protocol Suite’ paper). Give detailed steps showing how an attacker might predict the sequence numbers of a server. Assume the following–1) the server increments the sequence number (SN) every millisecond by 1; 2) at the start (boot) time,  $t = 0$  and  $SN = 0$ ; 3) the time window is 1ms, implying at  $t = 1$ ,  $SN = 1$ , at  $t = 2$ ,  $SN = 2$ , at  $t = 3$ ,  $SN = 3$ , and so on ( because the server increments SN every millisecond); and 4) the round-trip time between the attacker’s computer and the server is exactly 1 second. Feel free to make reasonable assumptions to answer this question. State your assumptions.**

# Attack on RIP

